

```
In [93]: #importing the libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [94]: #reading the dataset
data=pd.read_csv('car data.csv')
print(data.shape)
data.head()
```

(301, 9)

```
Out[94]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	Manual
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	Manual
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	Manual
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	Manual
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	Manual

From the dataset preview, Selling\_Price is the dependent variable and the rest are independent variables

```
In [95]: print("Top 10 cars:\n")
print(data.Car_Name.value_counts().head(10))

y = data.Car_Name.value_counts(normalize=True).cumsum()
x = len(y)

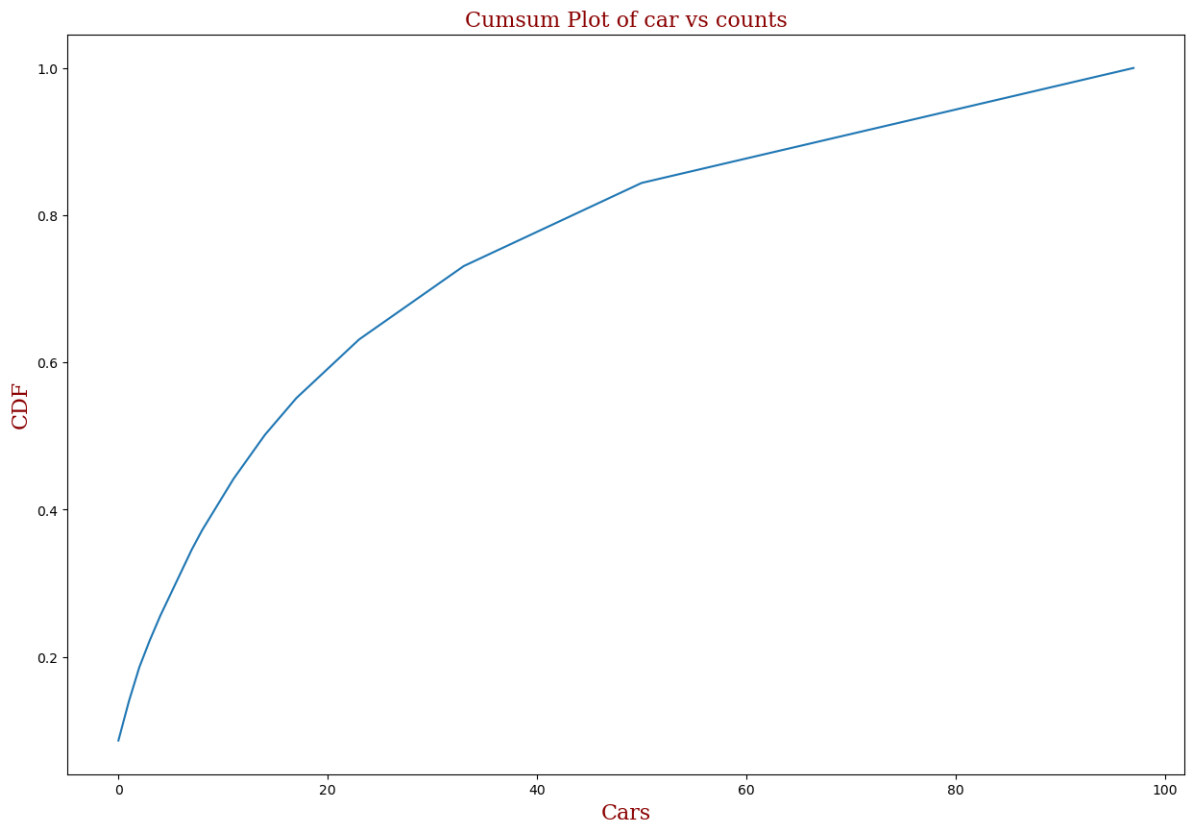
font = {'family': 'serif',
        'color': 'darkred',
        'weight': 'normal',
        'size': 16,
        }

plt.figure(figsize=(15,10))
plt.plot(range(x),y);
plt.title('Cumsum Plot of car vs counts', fontdict = font)
plt.xlabel('Cars', fontdict = font)
plt.ylabel('CDF', fontdict = font)
```

Top 10 cars:

```
city          26
corolla altis 16
verna         14
fortuner      11
brio          10
ciaz          9
innova        9
i20           9
grand i10     8
jazz          7
Name: Car_Name, dtype: int64
```

```
Out[95]: Text(0, 0.5, 'CDF')
```



## Checking if there are any missing values

In [96]: `data.isnull().sum()`

```
Out[96]: Car_Name      0
Year          0
Selling_Price 0
Present_Price 0
Kms_Driven    0
Fuel_Type     0
Seller_Type   0
Transmission  0
Owner         0
dtype: int64
```

## Checking cardinality of independent categorical variables in the dataset

```
In [97]: print('Unique elements in Seller_Type are',data['Seller_Type'].unique())
print('Unique elements in Fuel_Type are',data['Fuel_Type'].unique())
print('Unique elements in Transmission are',data['Transmission'].unique())
print('Unique elements in Owner are',data['Owner'].unique())
print('Unique elements in Year are',data['Year'].unique())
```

```
Unique elements in Seller_Type are ['Dealer' 'Individual']
Unique elements in Fuel_Type are ['Petrol' 'Diesel' 'CNG']
Unique elements in Transmission are ['Manual' 'Automatic']
Unique elements in Owner are [0 1 3]
Unique elements in Year are [2014 2013 2017 2011 2018 2015 2016 2009 2010 2012 200
3 2008 2006 2005
2004 2007]
```

```
In [98]: print('Unique elements in Car_Name are',data['Car_Name'].nunique())
#98 unique elements
#so, rather than encoding it, we can just drop this column as it doesn't make sense
```

Unique elements in Car\_Name are 98

```
In [99]: data.describe()
```

```
Out[99]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Owner
<b>count</b>	301.000000	301.000000	301.000000	301.000000	301.000000
<b>mean</b>	2013.627907	4.661296	7.628472	36947.205980	0.043189
<b>std</b>	2.891554	5.082812	8.644115	38886.883882	0.247915
<b>min</b>	2003.000000	0.100000	0.320000	500.000000	0.000000
<b>25%</b>	2012.000000	0.900000	1.200000	15000.000000	0.000000
<b>50%</b>	2014.000000	3.600000	6.400000	32000.000000	0.000000
<b>75%</b>	2016.000000	6.000000	9.900000	48767.000000	0.000000
<b>max</b>	2018.000000	35.000000	92.600000	500000.000000	3.000000

## Feature Engineering

```
In [100... dataset=data[['Year','Selling_Price','Present_Price','Kms_Driven','Fuel_Type','Seller_Type','Transmission','Owner']]
dataset.head()
```

```
Out[100]:
```

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner
<b>0</b>	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0
<b>1</b>	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0
<b>2</b>	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0
<b>3</b>	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0
<b>4</b>	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0

defining a feature variable 'Present\_Year' which has all the element values as 2020. On subtracting 'Present\_Year' and 'Year', we can make another feature variable as 'Number\_of\_Years\_Old', which gives us idea about how old the car is.

```
In [101... dataset['Present_Year']=2020
dataset['Number_of_Years_Old']=dataset['Present_Year']- dataset['Year']
dataset.head()
```

Out[101]:

	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	P
0	2014	3.35	5.59	27000	Petrol	Dealer	Manual	0	
1	2013	4.75	9.54	43000	Diesel	Dealer	Manual	0	
2	2017	7.25	9.85	6900	Petrol	Dealer	Manual	0	
3	2011	2.85	4.15	5200	Petrol	Dealer	Manual	0	
4	2014	4.60	6.87	42450	Diesel	Dealer	Manual	0	

So, we can now safely drop 'Year' and 'Present\_Year' columns

In [102...]

```
dataset.drop(labels=['Year', 'Present_Year'],axis=1,inplace=True)
dataset.head()
```

Out[102]:

	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission	Owner	Number
0	3.35	5.59	27000	Petrol	Dealer	Manual	0	
1	4.75	9.54	43000	Diesel	Dealer	Manual	0	
2	7.25	9.85	6900	Petrol	Dealer	Manual	0	
3	2.85	4.15	5200	Petrol	Dealer	Manual	0	
4	4.60	6.87	42450	Diesel	Dealer	Manual	0	

## Encoding the Categorical Variables

In [103...]

```
#select categorical variables from then dataset, and then implement categorical encoding
Fuel_Type=dataset[['Fuel_Type']]
Fuel_Type=pd.get_dummies(Fuel_Type, drop_first=True)

Seller_Type=dataset[['Seller_Type']]
Seller_Type=pd.get_dummies(Seller_Type, drop_first=True)

Transmission=dataset[['Transmission']]
Transmission=pd.get_dummies(Transmission, drop_first=True)

dataset=pd.concat([dataset,Fuel_Type, Seller_Type, Transmission], axis=1)

dataset.drop(labels=['Fuel_Type', 'Seller_Type', 'Transmission'], axis=1, inplace=True)
dataset.head()
```

Out[103]:

	Selling_Price	Present_Price	Kms_Driven	Owner	Number_of_Years_Old	Fuel_Type_Diesel	Fuel_Type_Petrol
0	3.35	5.59	27000	0	6	0	1
1	4.75	9.54	43000	0	7	1	0
2	7.25	9.85	6900	0	3	0	1
3	2.85	4.15	5200	0	9	0	1
4	4.60	6.87	42450	0	6	1	0

to reduce the multicollinearity we dropped cng --Multicollinearity occurs when two or more independent variables in a data frame have a high correlation with one another in a regression model.

In [104... dataset.columns

Out[104]: Index(['Selling\_Price', 'Present\_Price', 'Kms\_Driven', 'Owner', 'Number\_of\_Years\_Old', 'Fuel\_Type\_Diesel', 'Fuel\_Type\_Petrol', 'Seller\_Type\_Individual', 'Transmission\_Manual'], dtype='object')

In [105... *# Dataset Correlation*  
dataset.corr()

Out[105]:


	Selling_Price	Present_Price	Kms_Driven	Owner	Number_of_Years_Old	Fuel_Type_Diesel	Fuel_Type_Petrol
<b>Selling_Price</b>	1.000000	0.878983	0.029187	-0.088344	-0.236141	0.552339	-0.540571
<b>Present_Price</b>	0.878983	1.000000	0.203647	0.008057	0.047584	0.473306	-0.465244
<b>Kms_Driven</b>	0.029187	0.203647	1.000000	0.089216	0.524342	0.172515	-0.172874
<b>Owner</b>	-0.088344	0.008057	0.089216	1.000000	0.182104	-0.053469	0.055687
<b>Number_of_Years_Old</b>	-0.236141	0.047584	0.524342	0.182104	1.000000	-0.064315	0.039896
<b>Fuel_Type_Diesel</b>	0.552339	0.473306	0.172515	-0.053469	-0.064315	1.000000	0.000394
<b>Fuel_Type_Petrol</b>	-0.540571	-0.465244	-0.172874	0.055687	0.039896	0.000394	1.000000
<b>Seller_Type_Individual</b>	-0.550724	-0.512030	-0.101419	0.124269	0.039896	0.000394	0.000000
<b>Transmission_Manual</b>	-0.367128	-0.348715	-0.162510	-0.050316	-0.000394	0.000000	0.000000

## Data Visualization and Correlation

In [106... *#Correlations of features in dataset*  
corrmat = data.corr()  
top\_corr\_features = corrmat.index  
plt.figure(figsize=(10,10))


```
#Plot heat map
sns.heatmap(data[top_corr_features].corr(),annot=True,cmap="RdYlGn")
```

Out[106]: <AxesSubplot:>



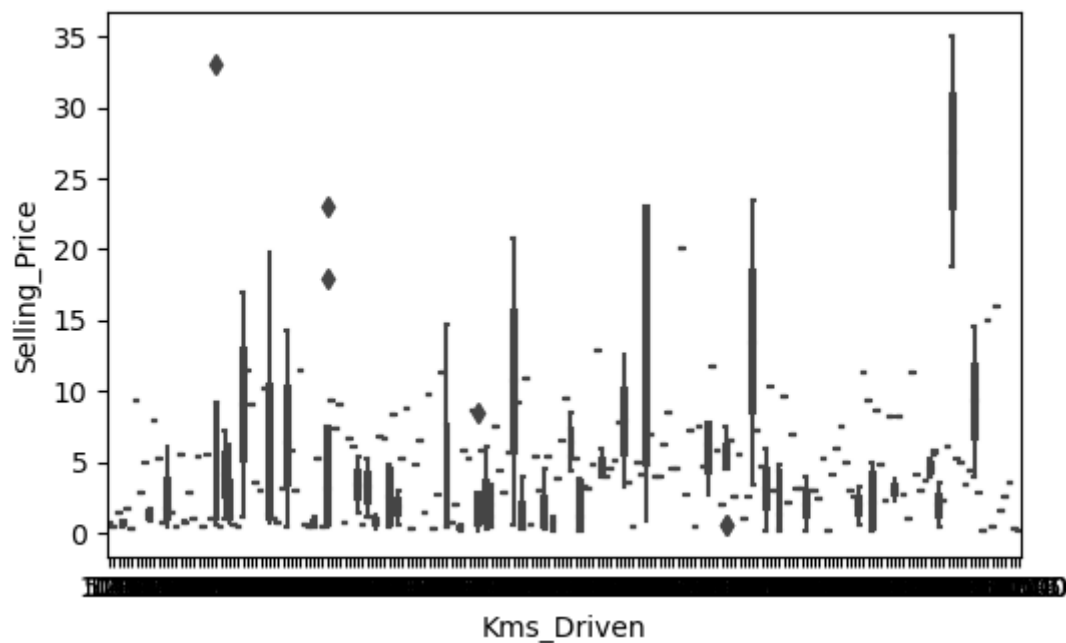
```
In [107... plt.figure(figsize = (20,12))
plt.subplot(3,3,1)
sns.boxplot(x = 'Fuel_Type', y = 'Selling_Price', data = data)
```

Out[107]: <AxesSubplot:xlabel='Fuel\_Type', ylabel='Selling\_Price'>

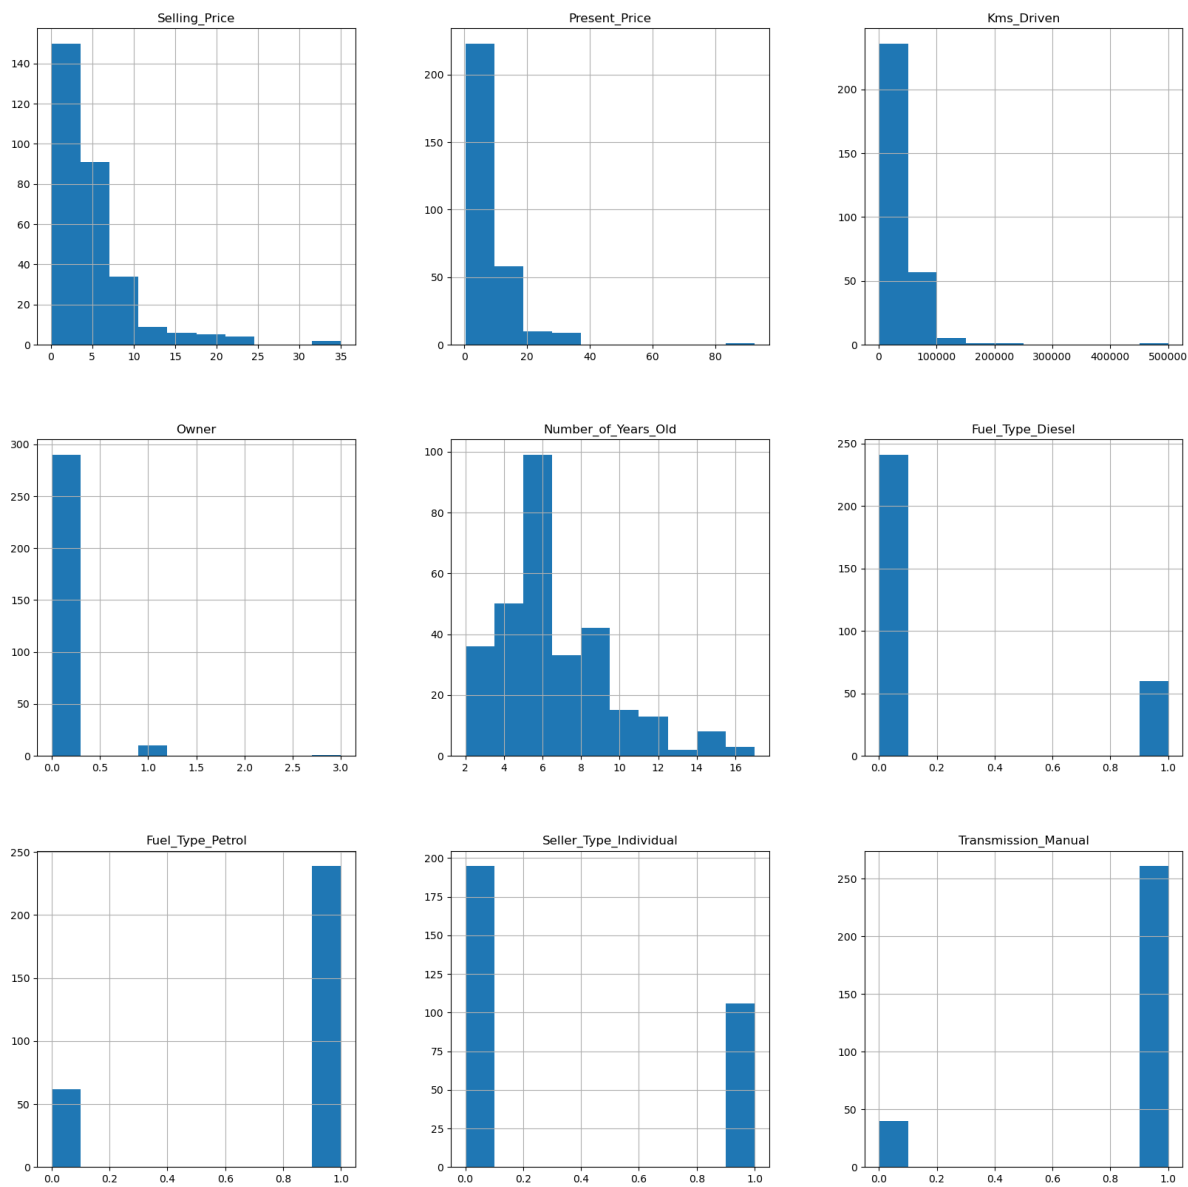


```
In [108... plt.figure(figsize = (20,12))
plt.subplot(3,3,1)
sns.boxplot(x = 'Kms_Driven', y = 'Selling_Price', data = data)
```

```
Out[108]: <AxesSubplot:xlabel='Kms_Driven', ylabel='Selling_Price'>
```




```
In [109... p = dataset.hist(figsize = (20,20))
```



```
In [110]: g=sns.PairGrid(dataset)
          g.map(plt.scatter)
```

```
Out[110]: <seaborn.axisgrid.PairGrid at 0x1844965cd30>
```






In [111...

```
def categorical_visualization(cols):
    plt.figure(figsize=(20,8))
    plt.subplot(1,3,1)
    sns.countplot(x=cols,data=data,palette="Set2",order=data[cols].value_counts().
    plt.title(f"{cols} Distribution",pad=10,fontweight="black",fontsize=18)
    plt.xticks(rotation=90)

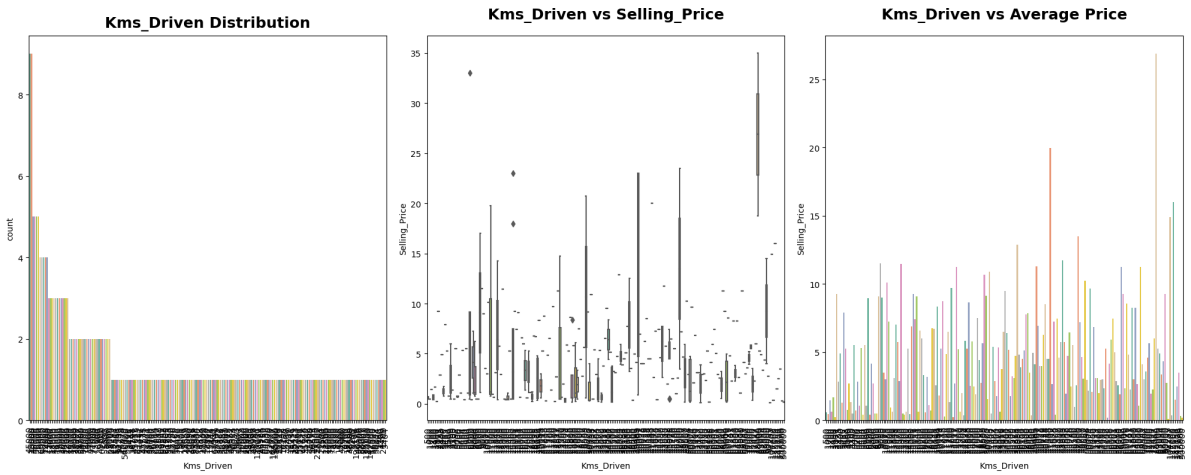
    plt.subplot(1,3,2)
    sns.boxplot(x=cols,y="Selling_Price",data=data,palette="Set2")
    plt.title(f"{cols} vs Selling_Price",pad=20,fontweight="black",fontsize=18)
    plt.xticks(rotation=90)

    plt.subplot(1,3,3)
    x=pd.DataFrame(data.groupby(cols)["Selling_Price"].mean().sort_values(ascending
    sns.barplot(x=x.index,y="Selling_Price",data=x,palette="Set2")
    plt.title(f"{cols} vs Average Price",pad=20,fontweight="black",fontsize=18)
    plt.xticks(rotation=90)
    plt.tight_layout()
    plt.show()

categorical_visualization("Fuel_Type")
```



```
In [112... categorical_visualization("Kms_Driven")
```



```
In [113... sell=dataset['Selling_Price']
dataset.drop(['Selling_Price'], axis=1, inplace=True)
dataset=dataset.join(sell)
dataset.head()
```

Out[113]:

	Present_Price	Kms_Driven	Owner	Number_of_Years_Old	Fuel_Type_Diesel	Fuel_Type_Petrol	S
0	5.59	27000	0	6	0	1	
1	9.54	43000	0	7	1	0	
2	9.85	6900	0	3	0	1	
3	4.15	5200	0	9	0	1	
4	6.87	42450	0	6	1	0	

```
In [114... X=dataset.iloc[:, :-1]
y=dataset.iloc[:, -1]
```


```
In [115... ### To determine important features, make use of ExtraTreesRegressor
from sklearn.ensemble import ExtraTreesRegressor
model = ExtraTreesRegressor()
model.fit(X,y)

print(model.feature_importances_)

#plot graph of feature importances for better visualization
feat_importances = pd.Series(model.feature_importances_, index=X.columns)
```

```
feat_importances_.nlargest(10).plot(kind='barh')
plt.show()
```

```
[0.4046579  0.04194815 0.0008649  0.07983709 0.21022608 0.01428844
 0.12603296 0.12214448]
```



'Owner' has zero feature importance i.e. nil on the dependent variable, 'Selling\_Price'

## Model Building and Training

```
In [116... X=dataset.iloc[:, :-1].values
y=dataset.iloc[:, -1].values
```

```
In [117... from sklearn.model_selection import cross_val_score
from sklearn import metrics
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
#from sklearn.model_selection import RandomizedSearchCV
#from sklearn.model_selection import GridSearchCV
#from sklearn.model_selection import StratifiedKFold
#kfold = StratifiedKFold(n_splits=3)
```

```
In [118... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [119... accuracy = []
r2score = []
mean2=[]
meanabsoulte=[]
rmse=[]
```

## Linear Regressor

```
In [120... from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

```

model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Linear Regressor Score on Training set is",model.score(X_train, y_train))#Training
print("Linear Regressor Score on Test Set is",model.score(X_test, y_test))#Testing

accuracies = cross_val_score(model, X_train, y_train, cv = 5)
print(accuracies)
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))

mae=mean_absolute_error(y_pred, y_test)
print("Mean Absolute Error:" , mae)

mse=mean_squared_error(y_test, y_pred)
print("Mean Squared Error:" , mse)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

print('The r2_score is', metrics.r2_score(y_test, y_pred))
accuracy.append(accuracies.mean()*100)
r2score.append(metrics.r2_score(y_test, y_pred))
mean2.append(mse)
meanabsoulte.append(mae)
rmse.append(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

sns.distplot(y_test-y_pred)
plt.show()

plt.scatter(y_test, y_pred, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()

```

```

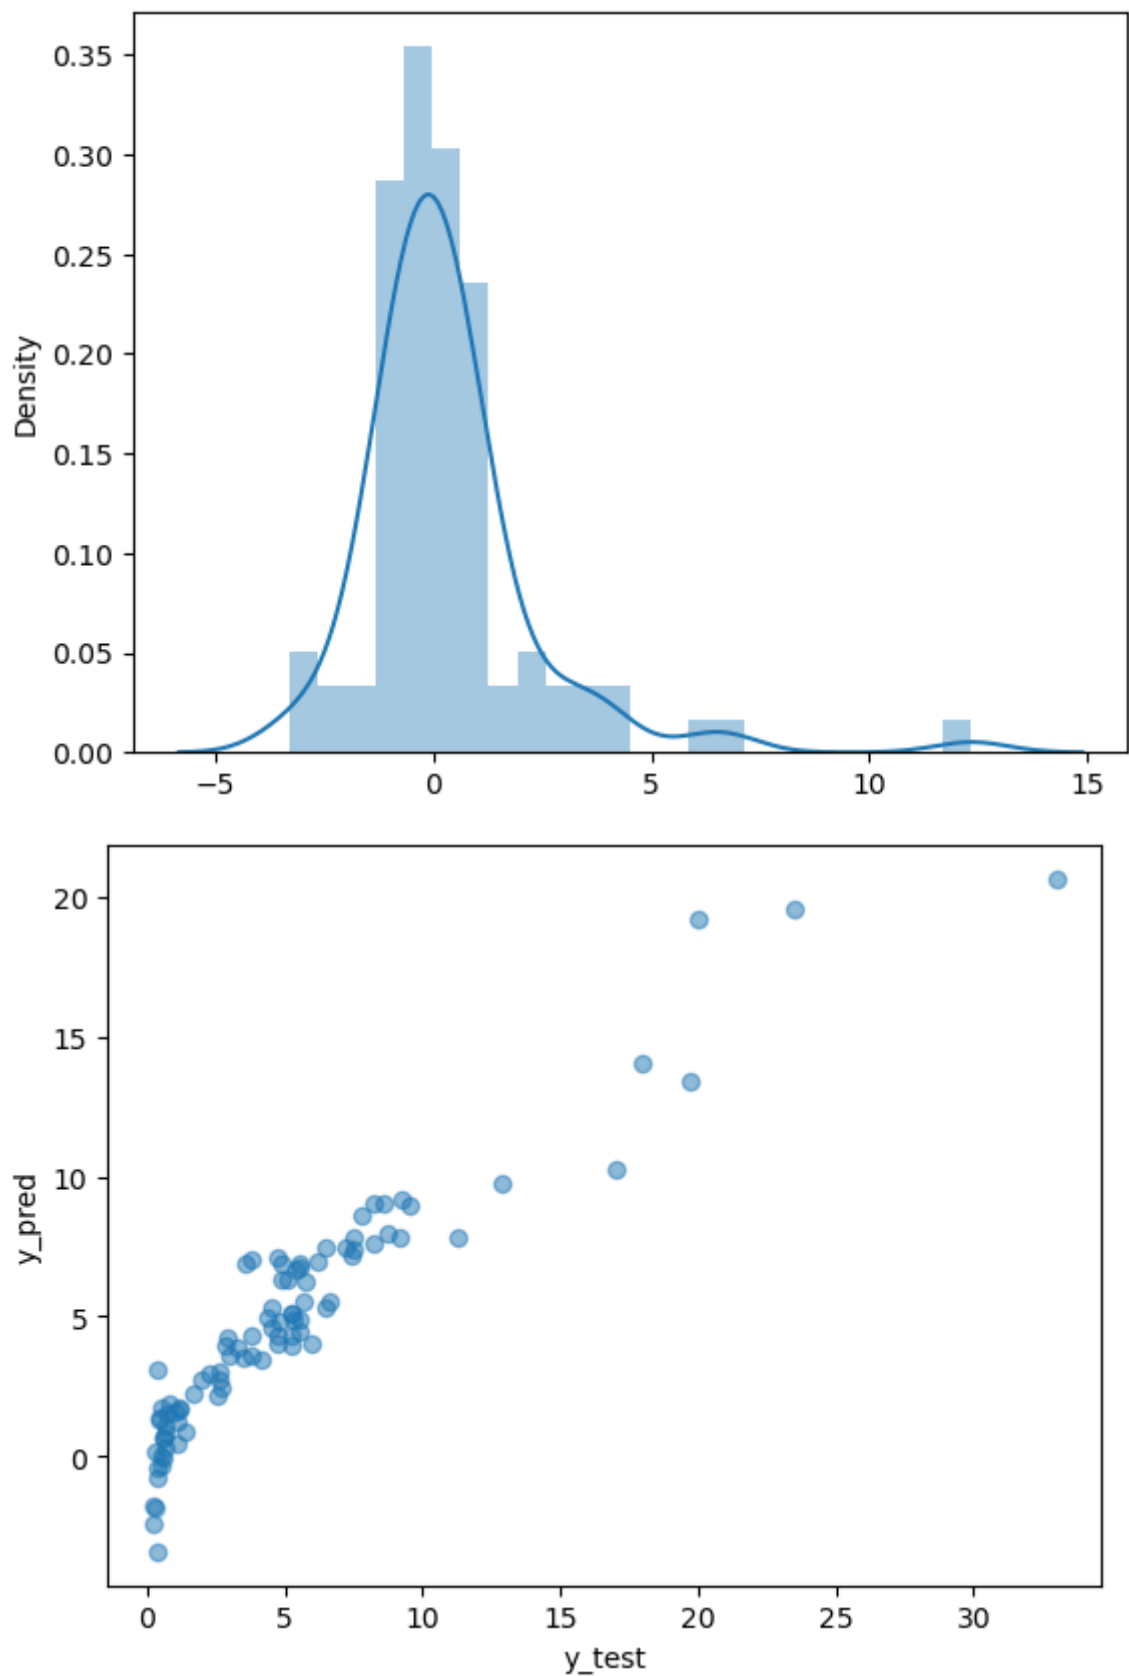
Linear Regressor Score on Training set is 0.8926207626208046
Linear Regressor Score on Test Set is 0.8517983059778262
[0.84989767 0.91045493 0.75641733 0.66228803 0.84093716]
Accuracy: 80.40 %
Standard Deviation: 8.62 %
Mean Absolute Error: 1.2426713915033707
Mean Squared Error: 4.432128265667619
RMSE: 2.1052620420431323
The r2_score is 0.8517983059778262

```


```

C:\Users\DELL\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)

```



```
In [121... sns.regplot(x=y_pred, y=y_test)
plt.xlabel("Predicted Price")
plt.ylabel('Actual Price')
plt.title("Actual vs predicted price")
plt.show()
```



## Decision Tree Regressor

```
In [122... #Decision Tree Regressor
from sklearn.tree import DecisionTreeRegressor
dt_reg = DecisionTreeRegressor(random_state = 0)
dt_reg.fit(X_train, y_train)
y_pred=dt_reg.predict(X_test)

print("Decision Tree Score on Training set is",dt_reg.score(X_train, y_train))#Training Score
print("Decision Tree Score on Test Set is",dt_reg.score(X_test, y_test))#Testing Accuracy

accuracies = cross_val_score(dt_reg, X_train, y_train, cv = 5)
print(accuracies)
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))

mae=mean_absolute_error(y_pred, y_test)
print("Mean Absolute Error:" , mae)

mse=mean_squared_error(y_test, y_pred)
print("Mean Squared Error:" , mse)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

print('The r2_score is', metrics.r2_score(y_test, y_pred))

accuracy.append(accuracies.mean()*100)
r2score.append(metrics.r2_score(y_test, y_pred))
mean2.append(mse)
meanabsoulte.append(mae)
rmse.append(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

sns.distplot(y_test-y_pred)
```

```
plt.show()

plt.scatter(y_test, y_pred, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()
```

Decision Tree Score on Training set is 1.0

Decision Tree Score on Test Set is 0.9202815383374512

[0.9542394 0.84409548 0.69916028 0.924205 0.92156403]

Accuracy: 86.87 %

Standard Deviation: 9.22 %


Mean Absolute Error: 0.8102197802197801


Mean Squared Error: 2.3840648351648355

RMSE: 1.5440417206684653

The r2\_score is 0.9202815383374512

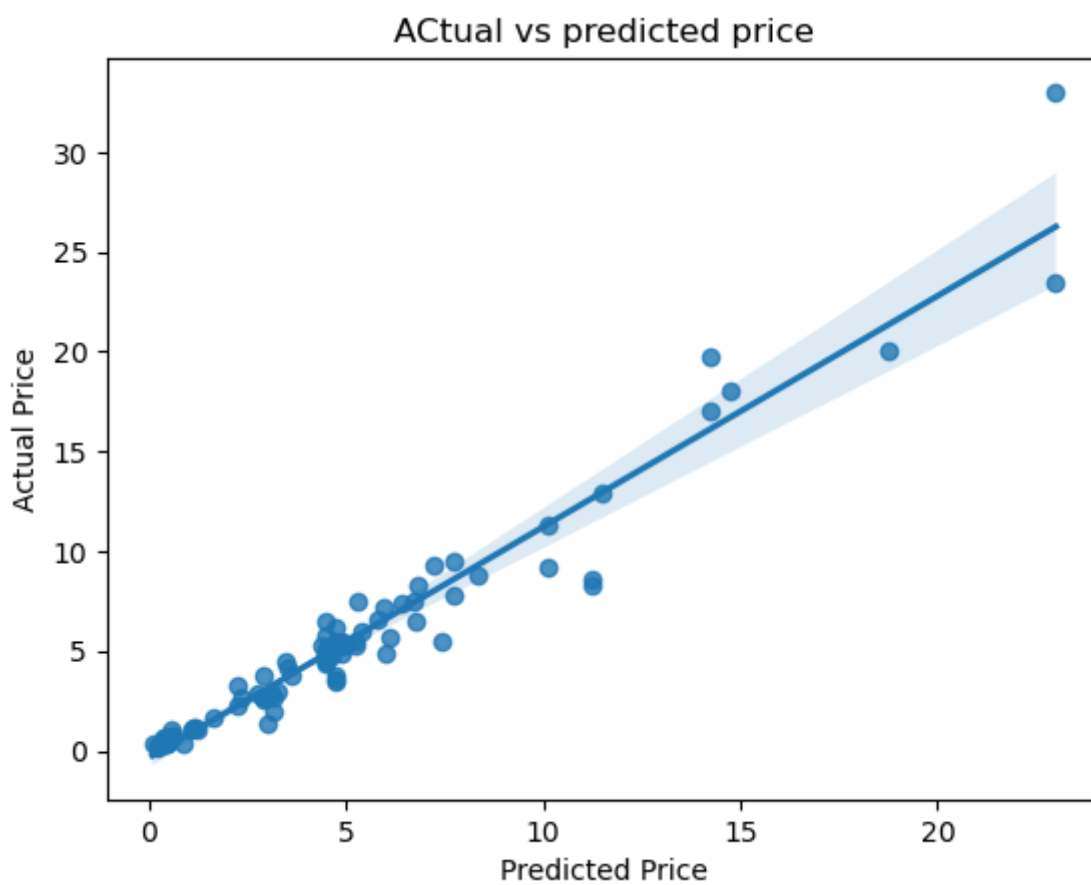
C:\Users\DELL\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)





In [123...

```
sns.regplot(x=y_pred, y=y_test)  
plt.xlabel("Predicted Price")  
plt.ylabel('Actual Price')  
plt.title("ACTual vs predicted price")  
plt.show()
```





# Random Forest Regressor

In [124...

```
#Random Forest Regression
from sklearn.ensemble import RandomForestRegressor
rf_reg = RandomForestRegressor(n_estimators=400,min_samples_split=15,min_samples_leaf=10,max_features='auto', max_depth=30)
rf_reg.fit(X_train, y_train)
y_pred=rf_reg.predict(X_test)

print("Random Forest Score on Training set is",rf_reg.score(X_train, y_train))#Training Accuracy
print("Random Forest Score on Test Set is",rf_reg.score(X_test, y_test))#Testing Accuracy

accuracies = cross_val_score(rf_reg, X_train, y_train, cv = 5)
print(accuracies)
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))

mae=mean_absolute_error(y_pred, y_test)
print("Mean Absolute Error:" , mae)

mse=mean_squared_error(y_test, y_pred)
print("Mean Squared Error:" , mse)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

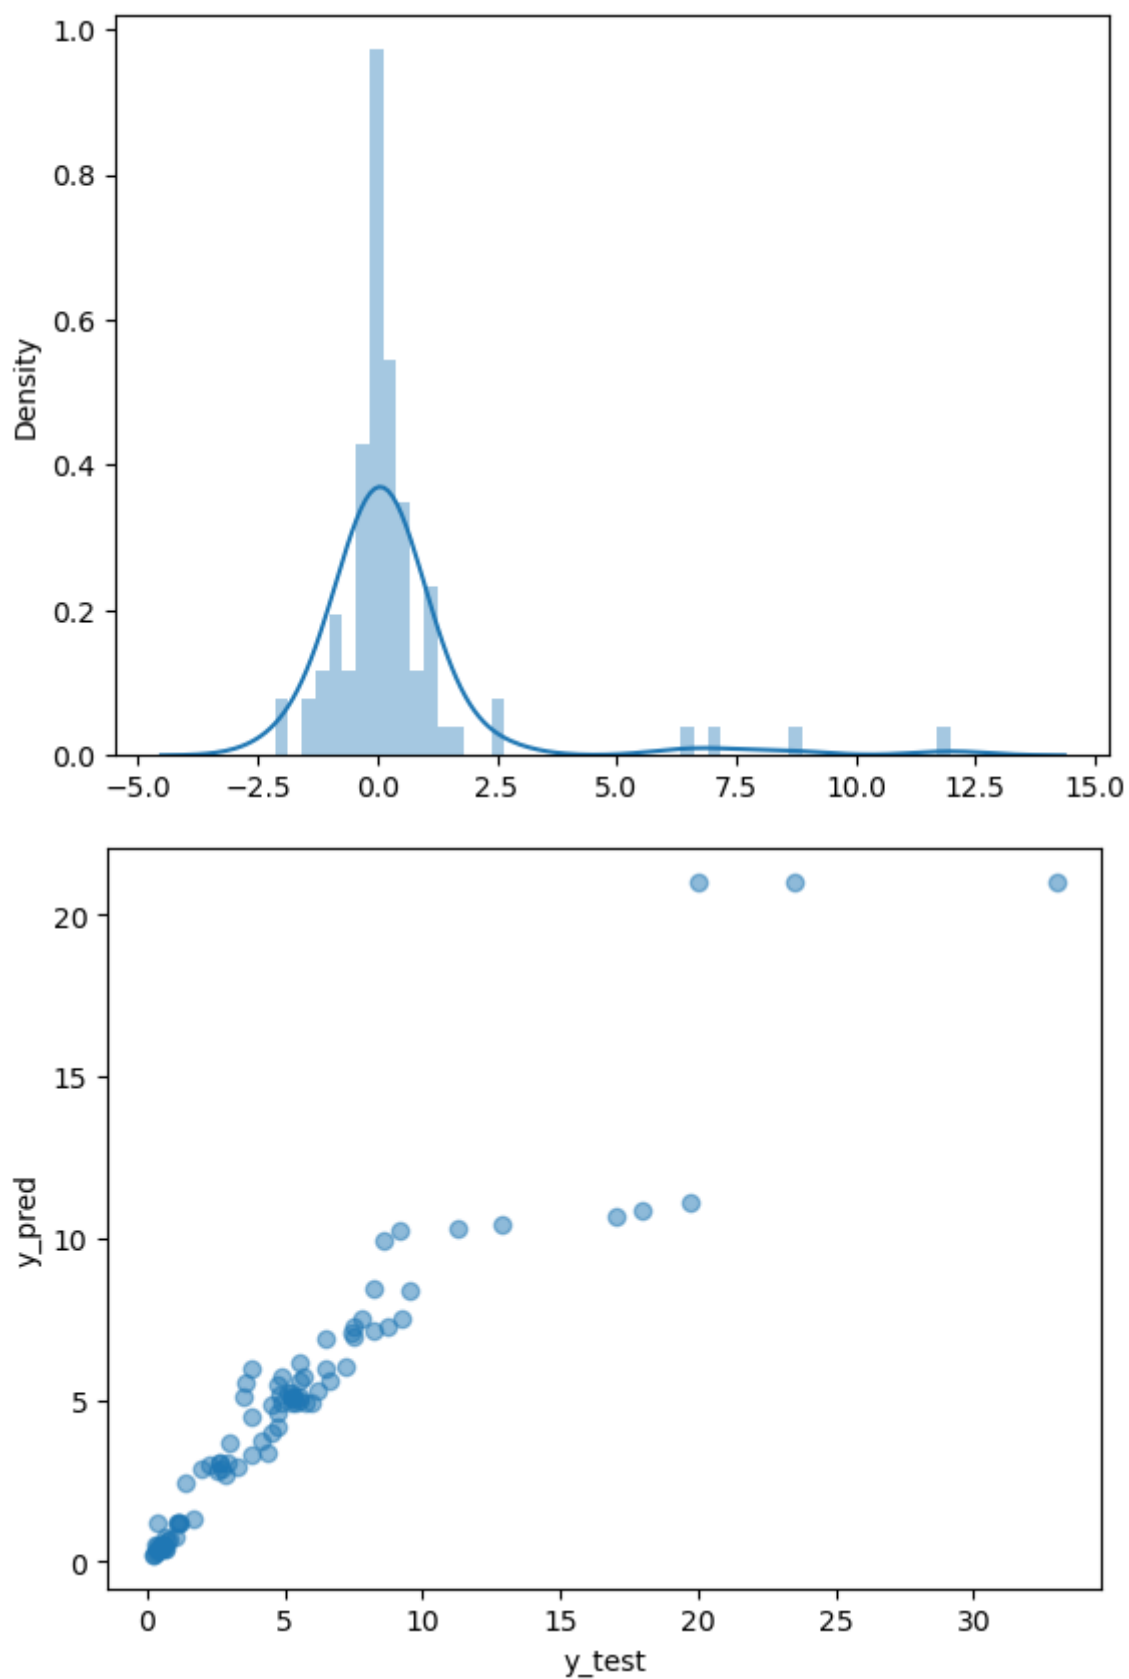
print('The r2_score is', metrics.r2_score(y_test, y_pred))
accuracy.append(accuracies.mean()*100)
r2score.append(metrics.r2_score(y_test, y_pred))
mean2.append(mse)
meanabsoulte.append(mae)
rmse.append(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

sns.distplot(y_test-y_pred)
plt.show()


plt.scatter(y_test, y_pred, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()
```

```
Random Forest Score on Training set is 0.910266568514969
Random Forest Score on Test Set is 0.8671354048467697
[0.94820538 0.816767 0.62988664 0.83125856 0.93028691]
Accuracy: 83.13 %
Standard Deviation: 11.34 %
Mean Absolute Error: 0.8836663992663134
Mean Squared Error: 3.973456117154846
RMSE: 1.993352983581896
The r2_score is 0.8671354048467697
```

```
C:\Users\DELL\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```



```
In [125... sns.regplot(x=y_pred, y=y_test)
plt.xlabel("Predicted Price")
plt.ylabel('Actual Price')
plt.title("ACTual vs predicted price")
plt.show()
```



## Voting Regressor

Voting Regressor is an ensemble meta-estimator that fits several base regressors, each on the whole dataset to average the individual predictions to form a final prediction.

In [126...

```
from sklearn.ensemble import VotingRegressor
vot_reg = VotingRegressor([('DecisionTree', dt_reg), ('RandomForestRegressor', rf_
vot_reg.fit(X_train, y_train)
y_pred=vot_reg.predict(X_test)

print("Voting Regressor Score on Training set is",vot_reg.score(X_train, y_train)
print("Voting Regressor Score on Test Set is",vot_reg.score(X_test, y_test))#Test

accuracies = cross_val_score(vot_reg, X_train, y_train, cv = 5)
print(accuracies)
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))

mae=mean_absolute_error(y_pred, y_test)
print("Mean Absolute Error:" , mae)

mse=mean_squared_error(y_test, y_pred)
print("Mean Squared Error:" , mse)

print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

print('The r2_score is', metrics.r2_score(y_test, y_pred))

accuracy.append(accuracies.mean()*100)
r2score.append(metrics.r2_score(y_test, y_pred))
mean2.append(mse)
meanabsoulte.append(mae)
```

```
rmse.append(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))

sns.distplot(y_test-y_pred)
plt.show()

plt.scatter(y_test, y_pred, alpha = 0.5)
plt.xlabel("y_test")
plt.ylabel("y_pred")
plt.show()
```

Voting Regressor Score on Training set is 0.9747670501881848

Voting Regressor Score on Test Set is 0.8983795201953172

[0.95785591 0.93777113 0.73902928 0.89954763 0.95815186]

Accuracy: 89.85 %

Standard Deviation: 8.25 %

Mean Absolute Error: 0.8544876881826632


Mean Squared Error: 3.039067831745924


RMSE: 1.743292239340818

The r2\_score is 0.8983795201953172

C:\Users\DELL\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


warnings.warn(msg, FutureWarning)





**So, the XGBoost Regressor gave us best scores on the Training Dataset**

```
In [127... sns.regplot(x=y_pred, y=y_test)
plt.xlabel("Predicted Price")
plt.ylabel('Actual Price')
plt.title("ACTual vs predicted price")
plt.show()
```



In [128... `models = ["Linear Regression","Decision Tree","Random Forest","X_boost"]`

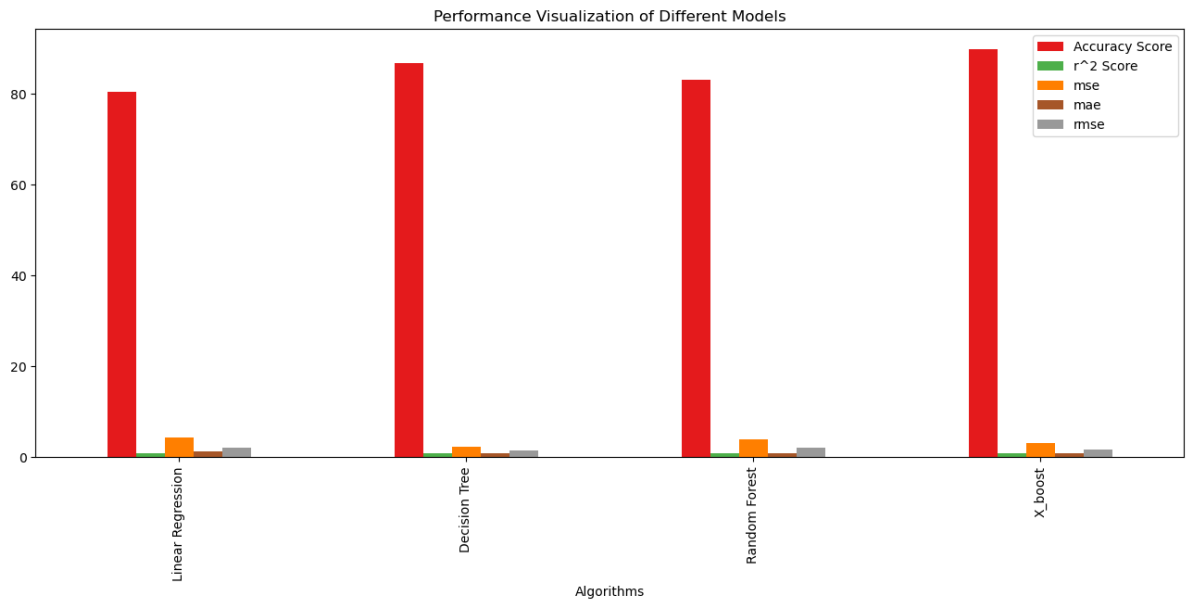
In [129... `df = pd.DataFrame({"Algorithms":models,  
"Accuracy Score":accuracy,  
"r^2 Score":r2score,  
"mse":mean2,  
"mae":meanabsoulte,  
"rmse":rmse})`

In [130... `df`

Out[130]:

	Algorithms	Accuracy Score	r^2 Score	mse	mae	rmse
0	Linear Regression	80.399903	0.851798	4.432128	1.242671	2.105262
1	Decision Tree	86.865284	0.920282	2.384065	0.810220	1.544042
2	Random Forest	83.128090	0.867135	3.973456	0.883666	1.993353
3	X_boost	89.847116	0.898380	3.039068	0.854488	1.743292

In [135... `df.plot(x="Algorithms",y=["Accuracy Score","r^2 Score","mse","mae","rmse"], figsize=  
title="Performance Visualization of Different Models",colormap="Set1")  
plt.show()`



## Model deployment by Dumping the model selected as a Pickle File

In [ ]:

```
import pickle
pickle.dump(vot_reg, open("vot_reg.pkl", "wb"))

# load model from file
model1 = pickle.load(open("vot_reg.pkl", "rb"))
```

```
In [132...] model1.predict([[9.85, 6900, 0, 3, 0, 1, 0, 1]])
```

Out[132]: array([7.19241174])

```
In [133...] model.predict([[7.85, 8000, 0, 2, 1, 2, 0, 1]])
```

Out[133]: array([9.69233381])

```
In [134...] model.predict([[11.0, 2000, 0, 1, 1, 2, 0, 1]])
```

Out[134]: array([11.34874994])

The car has more Selling price when the kms driven is less